

ByShard: Sharding in a Byzantine Environment

*Ramesh Adhikari, Graduate Research Assistant
School of Computer and Cyber Sciences, Augusta University*

September, 2022



AUGUSTA UNIVERSITY

Outline

- Motivation for this paper
- Used Protocol
- Main Idea
- Proposed Model
- Evaluation

Motivation for this paper

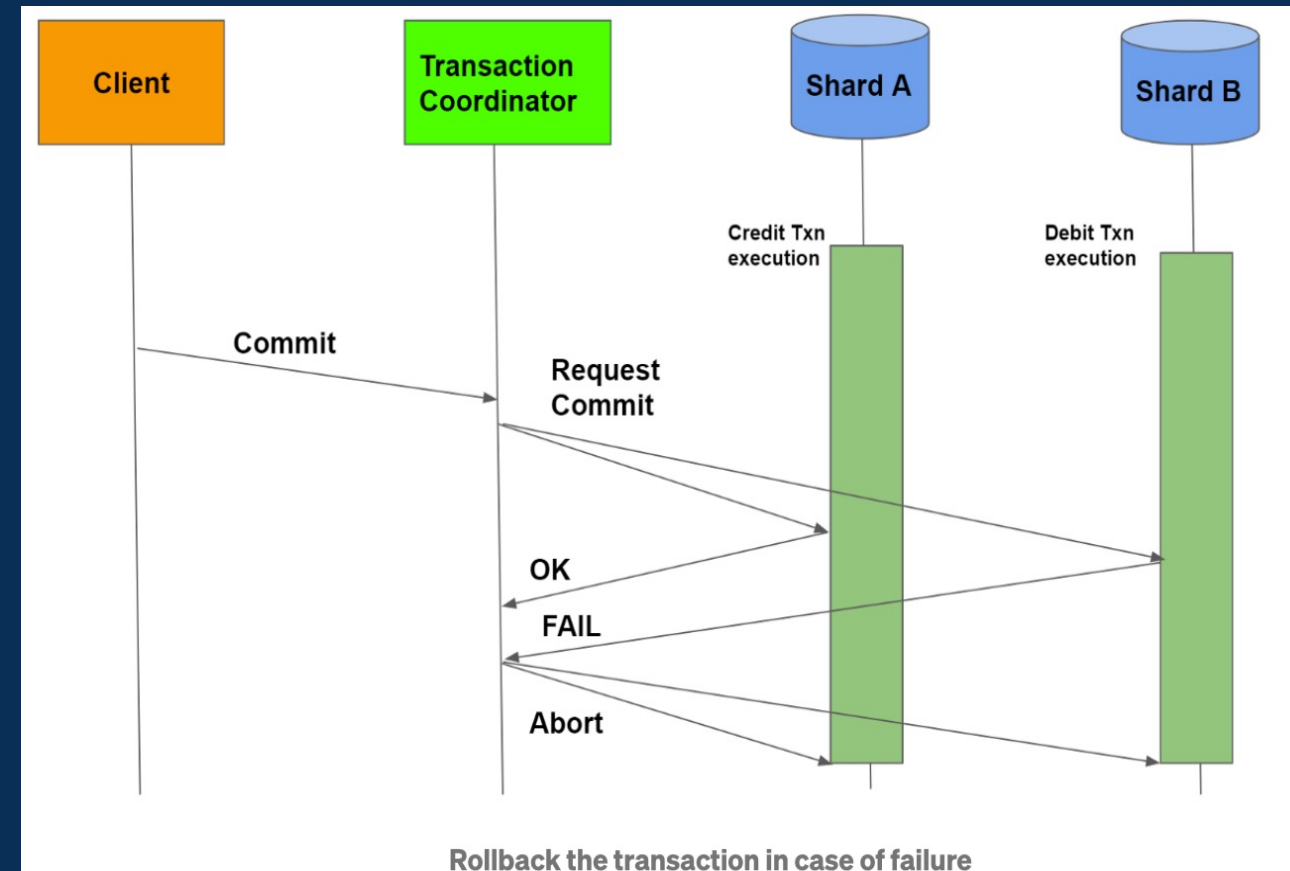

- Development of resilient system than can handle Byzantine failure due to Crashes, Bugs, Malicious behaviors
- Current Sharded resilient system do not provide the flexibility of traditional data management system
- To proposed High-Performance Resilient system

Used Protocol

Used two traditional sharded database concept efficiently in Byzantine environment

- **Two-phase commit:** Atomicity; atomic decision on whether the transaction can be committed or not;
- **Two-phase locking:** Isolation; provide concurrency control


Two Phase Commit



Two Phase Locking (Serializability)

Time	T1	T2
T ₀	Write Lock for A	
T ₁	Write Lock for B	
T ₂	Update A=A+1	
T ₃	Update B=B+2	
T ₄	Unlock A	
T ₅	Unlock B	
T ₆		Write Lock on A
T ₇		Write Lock on B
T ₈		Update A=A*2
T ₉		Update B=B*4
T ₁₀		Unlock A
T ₁₁		Unlock B

Sharded Design



Main Idea

- Orchestrate-execution model (OEM) in Byzantine environment
 - Orchestration: Replication of transactions among all involved shards and reaching on atomic decision; used two-phase commit
 - Execution model: Execution of transactions by maintaining data consistency among shards; used two-phase locking
- Uses cluster-sending communication
 - Particular algorithm unspecified
- Uses consensus abstraction as a Blackbox

Communication in Shard

- **Cluster-sending** protocol is used for reliable communication between clusters S_1 and S_2 ; To send S_1 value v to S_2 , Provide the following guarantees
 - S_1 can send v to S_2 only if there is agreement on sending v among the non-faulty replicas in S_1 ;
 - all non-faulty replicas in S_2 will receive the value v ; and
 - all non-faulty replicas in S_1 obtain confirmation of receipt.

Orchestrate-execution model (OEM)

- Processing is broken down into three types of shard-steps
 - **Vote-step:** Shard (S) Verifies the constraints to determine whether S votes for either commit or abort. And can make local changes, e.g., check conditions, modify local data or acquire locks
 - **Commit-step:** Shard performs necessary operations to finalize transactions when transactions is committed. E.g., Modify data and release locks
 - **Abort-step:** Shard performs necessary operations to roll-back transactions when transactions is aborted. E.g., roll -back local changes, release locks


Orchestration

- The main goal of it is to replicate the transactions (Tx) to involved shard and obtained the commit/abort decision
- Three type of model
 - Linear (based on Linear 2PC)
 - Centralized (based on 2PC)
 - Decentralized




Orchestration - Linear

- Vote Step: Sequence
- Decide: Centralized
- Commit or Abort: Parallel
- Advantage: Early abort
- S^1, S^2, S^3 and S^4 are vote-steps
- S^2, S^5 and S^6 have commit steps
- Every dot represents a single consensus step
- Every arrow a single cluster sending step




Orchestration - Centralized

- Root/Coordinator is selected for each Tx independently
- Vote Step: Parallel
- Decide: Centralized
- Commit or Abort: Parallel
- Disadvantage: Wait for all message



Orchestration - Decentralized

- Vote Step: Parallel
- Decide: Decentralized
- Commit or Abort: Parallel
- Can be performed in 3 consecutive steps
- Vote aggregation is performed in a single step as well




Execution Model

- Execution part consider the isolation
 - The above orchestrations allow to read uncommitted data
 - Two-phase locking is proposed to cope with that
 - A Tx is split to Constraint and Modification steps

Execution-Isolation free execution

- If S has a condition – update is made in the vote step
- Abort steps are generated for all such modification
- If S has no condition – modifications are made in the commit step, no abort step needed
- Disadvantage- Dirty read are possible



Execution – Lock-based execution

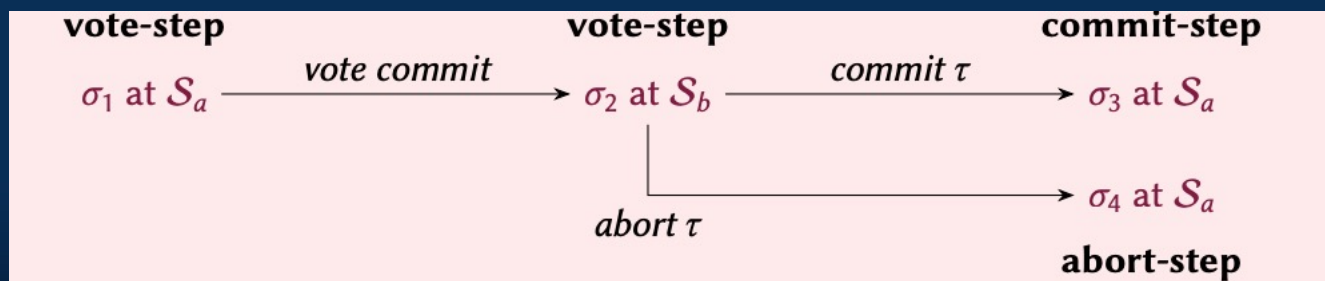
- Read/Write locks are used
- Modes:
 - Read uncommitted : Dirty Read
 - Read Committed : avoid a dirty read, but reads the same row twice and gets a different value each time
 - Serializable: read and write locks are used in a usual way; Two Phase Locking; data consistency; isolation

Example of the OEM

Shard accounts by first letter of name

Representations: τ is transaction; σ is shard-step;

- τ = “if Ana has \$500 and Bo has \$200, then move \$400 from Ana to Bo.
- σ_1 = “LOCK(**Ana**); if Ana has \$500, then forward σ_2 to Sb (Commit vote)
else RELEASE(**Ana**) (Abort vote)”
- σ_2 = “LOCK(**Bo**); if Bo has \$200, then add \$400 to Bo; RELEASE(**Bo**);
and forward σ_3 to Sa (Commit)
else RELEASE(**Bo**) and forward σ_4 to Sa (Abort)”
- σ_3 = “remove \$400 from Ana, Commit τ and RELEASE(**Ana**)”
- σ_4 = “Abort τ and RELEASE(**Ana**)”



Evaluation

- Consensus steps were abstracted in evaluations
- Experiment done on 5000 Tx
- Tx affects 16 accounts, 8 accounts have constrained
- 64 Shards
- 8k accounts
- Scalability increases with number shards while keeping other parameters constant.

Thank you!

